

Spatial Atlas discipline, pointed at **ad adjacency**

Typed scene state becomes a typed PolicySpec · Fail-closed execution becomes deterministic
verdict gates · Confidence-gated revision becomes selective reasoning

Adjacency

SYSTEM

Brand safety


SURFACE

Read-only

MODE

Keywords are legible, but the risk is often **between modalities**

- X exposes advertiser-controlled adjacency keyword exclusions
- A safe discussion can contain a blocked word
- A risky image can accompany clean text
- Source: src/adjacency/baseline.py and X Brand Safety documentation




Prose → terms

REPRESENTATION

Joint context

MISSING SIGNAL

Compile intent, inspect inventory, admit evidence, compare outcomes, route review



Source: ARCHITECTURE.md and src/adjacency/

- COMPILED, NOT PROMPTED

Every clause must point back to the words that authorized it

- PolicySpec is typed, immutable, versioned, hashed, and diffable
- GO checks exact character spans at compile time
- The keyword baseline is generated from the identical prose
- Sources: artifacts/tuesday/policy_spec.json and artifacts/tuesday/baseline_blocklist.json



GO

COMPILER ADMISSION

Same prose

BASELINE

Both directions matter: missed risk and recovered safe inventory

Adjacency Autopsy

Compare the compiled policy engine with the same-prose keyword baseline.

50 verified items · 18 local media

Source: `corpus/frozen/manifest.json`

Run Autopsy

Reset

GATE FAIL: G1_SPAN_NOT_FOUND

`trace:x-2054959102195872231` changed from **ALLOW** to **REVIEW** and entered human review.

AGREE

Both systems make the same delivery decision.

Agreement traces

Trace	Decision	Tier	Signal
<code>trace:x-2089017572486692982</code>	ALLOW → ALLOW	T0	Tier 0 deterministic clean-text allow.

OVER_BLOCK

The blacklist withholds safe inventory.

\$200.00

Illustrative recovered-spend counter · 2 Items × \$100.00

Source: `artifacts/ui/economics_assumption.json`

Recovered inventory traces

Trace	Decision	Tier	Signal
<code>trace:x-2082058501980045644</code>	BLOCK → ALLOW	T1	Matched: deepfakes, synthetic media

UNDER_BLOCK

The blacklist serves inventory the engine withholds.

GATE FAIL · G1_SPAN_NOT_FOUND · `trace:x-2054959102195872231` · ALLOW → REVIEW · Inspect

Caught-risk traces

Trace	Decision	Tier	Signal
<code>trace:x-2082911174522257413</code>	ALLOW → REVIEW	T2	GATE FAIL · G1_SPAN_NOT_FOUND

`trace:x-2054959102195872231`

UNDER_BLOCK · baseline **ALLOW** · engine **REVIEW** · Tier 2

Collection of 127 scientific diagrams for physics and ML <https://t.co/NlpJXOAP0h>

Evidence: Collection of 127 scientific diagrams for physics and ML, Title and subtitle confirming scientific diagrams for physics/chemistry/ML

[Open source post](#)

Audit drawer

Structured trace and full gate chain

```

1  {
2    "baseline": {
3      "action": "ALLOW",
4      "matched_terms": [
5        ]
6    }
7  }
```

Artifact-backed Autopsy after the G1 gate failure. Sources: artifacts/ui/autopsy_snapshot.json and artifacts/ui/economics_assumption.json. The dollar counter is illustrative.

A persuasive rationale cannot argue past **G1_SPAN_NOT_FOUND**

- The proposed ALLOW becomes REVIEW
- The row enters the human queue
- The audit drawer retains every applicable gate
- Source: artifacts/ui/autopsy_snapshot.json



ALLOW

BEFORE

REVIEW

AFTER

Fail closed


RULE

The strongest numbers test the failure path

OBSERVED RESULT	ARTIFACT
12/12 faults caught, 0/4 clean rejected	artifacts/tuesday/synthetic_faults.json
68% grounding failure, raw-prompt run 1	evals/prompt_baseline/comparison.json
2% cross-run action disagreement	evals/prompt_baseline/comparison.json
44% routed to Tier 2	artifacts/wednesday/judge_report.json
\$6.12 per 1,000 decisions	artifacts/wednesday/judge_report.json
Real Temporal round-trip completed	artifacts/temporal/hitl_live_proof.json

The measured ladder is **bimodal**, not a smooth curve

- Tier 0 absorbs the deterministic clean-text slice
- Tier 1 is a narrow residual in the frozen run
- Tier 2 handles the multimodal and uncertain cases
- Source: artifacts/wednesday/judge_report.json



Observed routing. Source: artifacts/wednesday/judge_report.json

44%

TIER 2 FRACTION

\$6.12 / 1,000

BLENDED JUDGE COST

Real Cloud review, offline demo, staged rollout

- Temporal proof: PENDING query → adjudication signal → ADJUDICATED query → completed workflow
- Artifact: artifacts/temporal/hitl_live_proof.json
- Demo default: in-process queue with no API key
- Rollout: read-only autopsy → shadow evaluation → per-clause activation

github.com/arunshar/adjacency

CODE

Models propose. Gates decide.

POSITION